

118430Q5_1

July 4, 2020

1 IA353 - Redes Neurais

2 EFC3 - Questão 5

2.1 Autoencoder training and manifold visualization

2.2 Partially based on <https://www.kaggle.com/apapiu/manifold-learning-and-autoencoders>

Professor: Fernando J. Von Zuben Aluno(a): Rafael Ito

2.2.1 5.1. Imports

```
[1]: import tensorflow as tf
import os
from multiprocessing import cpu_count

import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn import metrics
from sklearn.neighbors import NearestNeighbors

from keras.models import Sequential, Model
from keras.layers import Dense, Dropout, Convolution2D, MaxPooling2D, Flatten, Input, Conv2D, UpSampling2D, Reshape, Deconvolution2D, Conv2DTranspose, BatchNormalization
from keras.optimizers import adam
from keras.utils.np_utils import to_categorical
```

Using TensorFlow backend.

```
[2]: # version
tf.__version__
```

```
[2]: '2.2.0'
```

```
[3]: #-----
# additional config
#-----
```

```
# random seed generator
os.environ['PYTHONHASHSEED']=str(42)
np.random.seed(42)
tf.random.set_seed(42)
os.environ['TF_DETERMINISTIC_OPS'] = '1'
```

```
[ ]: # choose between CPU and GPU
device = tf.device('/cpu:0')
if tf.config.list_physical_devices('GPU'):
    device = tf.device('/device:GPU:0')
#     device_model = torch.cuda.get_device_name(0)
#     device_memory = torch.cuda.get_device_properties(device).total_memory / 1e9
device_number = len(tf.config.experimental.list_physical_devices('GPU'))
#-----
#from tensorflow.python.client import device_lib
#print(device_lib.list_local_devices())
#-----
print('Device: gpu')
#print('GPU model:', device_model)
#print('GPU memory: {0:.2f} GB'.format(device_memory))
print("GPUs available: ", device_number)
print('#-----')
print('CPU cores:', cpu_count())
```

2.2.2 5.2. Reading the data

```
[5]: mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```


Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
[6]: x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
[7]: x_images = x_train.reshape(x_train.shape[0], 28, 28)
```

```
[8]: # Definition of a function to visualize some digits
def show(img):
    plt.imshow(img, cmap = "gray", interpolation = "none")
```

```
[9]: # Visualization of 25 x 25 original images of digits
fig = plt.figure(figsize=(28,28))
ind = 1
for i in range(1, 26, 1):
    for j in range(1, 26, 1):
        fig.add_subplot(25,25,ind)
        show(x_images[ind])
        plt.xticks([])
        plt.yticks([])
        ind+=1
```



2.2.3 5.3 Proposal for the autoencoder architecture

```
[10]: input_img = Input(shape=(784,))
#=====
reshape = Reshape((28,28,1))(input_img)
encoded = Conv2D(32, (3, 3), activation='relu', strides=2, padding='same')(reshape)
→      # (60000,14,14,32)
encoded = Conv2D(64, (3, 3), activation='relu', strides=2, padding='same')(encoded)
→      # (60000,7,7,64)
encoded = Conv2D(128, (3, 3), activation='relu', strides=2, padding='valid')(encoded)
→      # (60000,3,3,128)
#-----
encoded = Flatten()(encoded)                                # (60000,1152)
encoded = Dense(128, activation='tanh')(encoded)            # (60000,128)
encoded = Dense(16, activation='tanh')(encoded)             # (60000,16)
```

```

#=====
encoded = Dense(2, activation=None)(encoded)           # (60000,2) --> bottleneck
#=====
decoded = Dense(16, activation='tanh')(encoded)        # (60000,16)
decoded = Dense(128, activation='tanh')(decoded)       # (60000,128)
decoded = Dense(1152, activation='tanh')(decoded)      # (60000,1152)
decoded = Reshape((3, 3, 128))(decoded)                # (60000,3,3,128)
→
#-----
decoded = Conv2DTranspose(64, (3,3), strides=2, activation='relu',
→padding='valid')(decoded)    # (60000,7,7,64)
decoded = BatchNormalization()(decoded)
#-----
decoded = Conv2DTranspose(32, (3,3), strides=2, activation='relu',
→padding='same')(decoded)    # (60000,14,14,32)
decoded = BatchNormalization()(decoded)
#-----
decoded = Conv2DTranspose(1, (3,3), strides=2, activation='sigmoid',
→padding='same')(decoded)    # (60000,28,28,1)
decoded = Reshape((784,))(decoded)
#=====
autoencoder = Model(input_img, output=decoded)

```

[12]: autoencoder.summary()

Model: "model_2"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 784)	0
reshape_1 (Reshape)	(None, 28, 28, 1)	0
conv2d_1 (Conv2D)	(None, 14, 14, 32)	320
conv2d_2 (Conv2D)	(None, 7, 7, 64)	18496
conv2d_3 (Conv2D)	(None, 3, 3, 128)	73856
flatten_1 (Flatten)	(None, 1152)	0
dense_1 (Dense)	(None, 128)	147584
dense_2 (Dense)	(None, 16)	2064
dense_3 (Dense)	(None, 2)	34
dense_4 (Dense)	(None, 16)	48
dense_5 (Dense)	(None, 128)	2176
dense_6 (Dense)	(None, 1152)	148608

reshape_2 (Reshape)	(None, 3, 3, 128)	0

conv2d_transpose_1 (Conv2DTr	(None, 7, 7, 64)	73792

batch_normalization_1 (Batch	(None, 7, 7, 64)	256

conv2d_transpose_2 (Conv2DTr	(None, 14, 14, 32)	18464

batch_normalization_2 (Batch	(None, 14, 14, 32)	128

conv2d_transpose_3 (Conv2DTr	(None, 28, 28, 1)	289


reshape_3 (Reshape)	(None, 784)	0
=====		
Total params: 486,115		
Trainable params: 485,923		
Non-trainable params: 192		

```
[14]: with device:
      autoencoder.compile(optimizer='adam', loss = "mse")
      X = x_train.reshape(x_train.shape[0], 784)
      history = autoencoder.fit(
          X,
          X,
          batch_size = 128,
          nb_epoch = 10,
          verbose = 1,
      )
```


```
Epoch 1/10
60000/60000 [=====] - 13s 222us/step - loss: 0.0718
Epoch 2/10
60000/60000 [=====] - 6s 105us/step - loss: 0.0454
Epoch 3/10
60000/60000 [=====] - 6s 103us/step - loss: 0.0430
Epoch 4/10
60000/60000 [=====] - 6s 102us/step - loss: 0.0417
Epoch 5/10
60000/60000 [=====] - 6s 104us/step - loss: 0.0407
Epoch 6/10
60000/60000 [=====] - 6s 103us/step - loss: 0.0399
Epoch 7/10
60000/60000 [=====] - 6s 104us/step - loss: 0.0393
Epoch 8/10
60000/60000 [=====] - 6s 103us/step - loss: 0.0388
Epoch 9/10
60000/60000 [=====] - 6s 102us/step - loss: 0.0383
Epoch 10/10
60000/60000 [=====] - 6s 102us/step - loss: 0.0380
```

```
[15]: plot_df = pd.DataFrame.from_dict({'train_loss':history.history['loss']})
      plot_df.plot(logy=True, figsize=(15,5), fontsize=12)
      plt.xlabel('epoch', fontsize=12)
```

```
plt.ylabel('loss', fontsize=12)
plt.show()
```



```
[16]: # Visualization of 25 x 25 reconstructed images of digits
fig = plt.figure(figsize=(28,28))
ind = 1
for i in range(1, 26, 1):
    for j in range(1, 26, 1):
        fig.add_subplot(25,25,ind)
        show(autoencoder.predict(np.expand_dims(x_images[ind].flatten(), 0)).
        ↪reshape(28, 28))
        plt.xticks([])
        plt.yticks([])
        ind+=1
```



```
[18]: encoder = Model(input = input_img, output = encoded)
```

```
# Building the decoder
encoded_input = Input(shape=(2,))
encoded_layer_1 = autoencoder.layers[-10]
encoded_layer_2 = autoencoder.layers[-9]
encoded_layer_3 = autoencoder.layers[-8]
encoded_layer_4 = autoencoder.layers[-7]
encoded_layer_5 = autoencoder.layers[-6]
encoded_layer_6 = autoencoder.layers[-5]
encoded_layer_7 = autoencoder.layers[-4]
encoded_layer_8 = autoencoder.layers[-3]
encoded_layer_9 = autoencoder.layers[-2]
encoded_layer_10 = autoencoder.layers[-1]

decoder = encoded_layer_1(encoded_input)
```

```

decoder = encoded_layer_2(decoder)
decoder = encoded_layer_3(decoder)
decoder = encoded_layer_4(decoder)
decoder = encoded_layer_5(decoder)
decoder = encoded_layer_6(decoder)
decoder = encoded_layer_7(decoder)
decoder = encoded_layer_8(decoder)
decoder = encoded_layer_9(decoder)
decoder = encoded_layer_10(decoder)
decoder = Model(input=encoded_input, output=decoder)

```

2.2.4 5.4 Visualizing the mapping of the labeled images in the manifold

```

[19]: import seaborn as sns

x_flat = x_train.reshape(x_train.shape[0], x_train.shape[1] * x_train.shape[2])

x_proj = encoder.predict(x_flat[:60000])
x_proj.shape


proj = pd.DataFrame(x_proj)
proj.columns = ["comp_1", "comp_2"]
proj["labels"] = y_train[:60000]
sns.lmplot("comp_1", "comp_2", hue="labels", data=proj, fit_reg=False)

```

```

[19]: <seaborn.axisgrid.FacetGrid at 0x7fb9da4d3ba8>

```





2.2.5 5.5. Generating new digits by moving along the manifold (latent 2D space)

Please, adjust the scale whenever necessary.

2.2.6 5.5.1 Moving along a circle

```
[38]: _ = sns.lmplot("comp_1", "comp_2", hue="labels", data=proj, fit_reg=False)
all_x = [2.5 * np.cos(2 * np.pi / 100 * i) + 0.5 for i in range(100)]
all_y = [2.5 * np.sin(2 * np.pi / 100 * i) + 0 for i in range(100)]
_ = plt.plot(all_x, all_y, 'black')
```




```
[39]: # moving along a circle:
_ = plt.figure(figsize=(10, 10))

for i in range(100):
    _ = plt.subplot(10, 10, i+1)


    x = 3 * np.cos(2.5 * np.pi / 100 * i) + 0.5
    y = 3 * np.sin(2.5 * np.pi / 100 * i) + 0

    pt = np.array([[x, y]])
    _ = show(decoder.predict(pt).reshape((28, 28)))
    _ = plt.xticks([])
    _ = plt.yticks([])
```



2.2.7 5.5.2 Moving along an ellipse

```
[42]: _ = sns.lmplot("comp_1", "comp_2", hue="labels", data=proj, fit_reg=False)
all_x = [1.5 * np.cos(2 * np.pi / 100 * i) - 0.5 for i in range(100)]
all_y = [2.5 * np.sin(2 * np.pi / 100 * i) - 0 for i in range(100)]
_ = plt.plot(all_x, all_y, 'black')
```




```
[43]: # moving along a ellipse:
_ = plt.figure(figsize=(10, 10))

for i in range(100):
    _ = plt.subplot(10, 10, i+1)

    x = 1.5 * np.cos(2 * np.pi / 100 * i) - 0.5
    y = 2.5 * np.sin(2 * np.pi / 100 * i) - 0

    pt = np.array([[x, y]])
    _ = show(decoder.predict(pt).reshape((28, 28)))
    _ = plt.xticks([])
    _ = plt.yticks([])
```



2.2.8 5.5.3 Moving along a grid

```
[45]: # moving along a ellipse:
_ = plt.figure(figsize=(12, 12))

for i in range(20):
    for j in range(20):

        _ = plt.subplot(20, 20, i * 20 + j + 1)
        # range adapted to go from -2 to +2 in both X and Y axis
        pt = np.array([[ -2 + 4 * i/19, -2 + 4 * j/19]])
        _ = show(decoder.predict(pt).reshape((28, 28)))
        _ = plt.xticks([])
        _ = plt.yticks([])
```

