

118430Q5_2

July 7, 2020

IA353 - Redes Neurais

EFC3 - Questão 5

DAE - Denoising Autoencoder (training and visualization)

dataset: CIFAR-10

Professor: Fernando J. Von Zuben **Aluno(a):** Rafael Ito

5.1. Imports

```
[ ]: import tensorflow as tf
import os
from multiprocessing import cpu_count

import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn import metrics
from sklearn.neighbors import NearestNeighbors

from keras.models import Sequential, Model
from keras.layers import Dense, Dropout, Convolution2D, MaxPooling2D, Flatten, Input,
↳Conv2D, UpSampling2D, Reshape, Deconvolution2D, Conv2DTranspose, BatchNormalization,
↳LeakyReLU, add
from keras.optimizers import adam
from keras.utils.np_utils import to_categorical
```

Using TensorFlow backend.

```
[ ]: # version
tf.__version__
```

```
[ ]: '2.2.0'
```

```
[ ]: #-----
# additional config
#-----
```

```
# random seed generator
os.environ['PYTHONHASHSEED']=str(42)
np.random.seed(42)
tf.random.set_seed(42)
os.environ['TF_DETERMINISTIC_OPS'] = '1'
```

```
[ ]: # choose between CPU and GPU
device = tf.device('/cpu:0')
if tf.config.list_physical_devices('GPU'):
    device = tf.device('/device:GPU:0')
#     device_model = torch.cuda.get_device_name(0)
#     device_memory = torch.cuda.get_device_properties(device).total_memory / 1e9
device_number = len(tf.config.experimental.list_physical_devices('GPU'))
#-----
#from tensorflow.python.client import device_lib
#print(device_lib.list_local_devices())
#-----
print('Device: gpu')
#print('GPU model:', device_model)
#print('GPU memory: {0:.2f} GB'.format(device_memory))
print("GPUs available: ", device_number)
print('#-----')
print('CPU cores:', cpu_count())
```

```
Device: gpu
GPUs available: 1
#-----
CPU cores: 2
```

5.2. Reading the data

```
[ ]: cifar10 = tf.keras.datasets.cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```
[ ]: x_train, x_test = x_train / 255.0, x_test / 255.0
```


```
[ ]: x_images = x_train.reshape(x_train.shape[0], 32, 32, 3)
```

```
[ ]: # Definition of a function to visualize some images
def show(img):
    plt.imshow(img, cmap = None, interpolation = "none")
```

```
[ ]: # Visualization of 25 x 25 original images
fig = plt.figure(figsize=(32,32))
ind = 1
for i in range(1, 26, 1):
    for j in range(1, 26, 1):
        fig.add_subplot(25,25,ind)
        show(x_images[ind])
        plt.xticks([])
        plt.yticks([])
        ind+=1
```

Original images

```
[ ]: # Visualization of 5 x 5 original images
fig = plt.figure(figsize=(32,32))
ind = 1
for i in range(1, 6, 1):
    for j in range(1, 6, 1):
        fig.add_subplot(5,5,ind)
        show(x_images[ind])
        plt.xticks([])
        plt.yticks([])
        ind+=1
```



Noisy images

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



5.3 Proposal for the autoencoder architecture

Architecture and training history:

- Step 1: Training autoencoder with images without noise, just aiming image reconstruction.
- Step 2: Added BatchNorm layers.
- Step 3: Added Dropout layers.
- Step 4: Added skip connection (shortcut).

Reconstruction OK!

- Step 5: Added noise to input images and trained again.
- Step 6: Parameters adjusted.
- Step 7: Final training (architecture below).

```
[ ]: input_img = Input(shape=(32,32,3,))
      ↳ # (50000,32,32,3)
      #=====
      encoded = Conv2D(32, (3, 3), activation='relu', strides=2, padding='same')(input_img)
      ↳ # (50000,32,32,32)
      encoded = BatchNormalization()(encoded)
      encoded = Dropout(0.5)(encoded)
      #-----
      skip = Conv2D(64, (3, 3), activation=None, strides=2, padding='same')(encoded)
      ↳ # (50000,16,16,32)
      encoded = LeakyReLU()(skip)
      encoded = BatchNormalization()(encoded)
      encoded = Dropout(0.5)(encoded)
      #-----
      encoded = Conv2D(128, (3, 3), activation='relu', strides=2, padding='same')(encoded)
      ↳ # (50000,8,8,64)
      encoded = BatchNormalization()(encoded)
      #-----
      encoded = Flatten()(encoded)
      # (50000,1024)
      #=====
      encoded = Dense(50, activation=None)(encoded)
      # (50000,10) --> bottleneck
      #=====
      decoded = Dense(2048, activation='tanh')(encoded)
      # (50000,1152)
      decoded = Reshape((4, 4, 128))(decoded)
      ↳ # (50000,4,4,128)
      #-----
      decoded = Conv2DTranspose(64, (3,3), strides=2, activation='relu',
      ↳padding='same')(decoded)
      # (50000,8,8,64)
      decoded = add([decoded, skip])
      decoded = LeakyReLU()(decoded)
      decoded = BatchNormalization()(decoded)
      #-----
      decoded = Conv2DTranspose(32, (3,3), strides=2, activation=None,
      ↳padding='same')(decoded)
      # (50000,16,16,32)
      decoded = BatchNormalization()(decoded)
      decoded = Dropout(0.5)(decoded)
      #-----
      decoded = Conv2DTranspose(3, (3,3), strides=2, activation='sigmoid',
      ↳padding='same')(decoded)
      # (50000,32,32,3)
      #=====
      autoencoder = Model(input=input_img, output=decoded)
```

```
[ ]: autoencoder.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 32, 32, 3)	0	
conv2d_1 (Conv2D)	(None, 16, 16, 32)	896	input_1[0][0]
batch_normalization_1 (BatchNormalizati	(None, 16, 16, 32)	128	conv2d_1[0][0]
dropout_1 (Dropout)	(None, 16, 16, 32)	0	batch_normalization_1[0][0]
conv2d_2 (Conv2D)	(None, 8, 8, 64)	18496	dropout_1[0][0]
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 64)	0	conv2d_2[0][0]
batch_normalization_2 (BatchNormalizati	(None, 8, 8, 64)	256	leaky_re_lu_1[0][0]
dropout_2 (Dropout)	(None, 8, 8, 64)	0	batch_normalization_2[0][0]
conv2d_3 (Conv2D)	(None, 4, 4, 128)	73856	dropout_2[0][0]
batch_normalization_3 (BatchNormalizati	(None, 4, 4, 128)	512	conv2d_3[0][0]
flatten_1 (Flatten)	(None, 2048)	0	batch_normalization_3[0][0]
dense_1 (Dense)	(None, 50)	102450	flatten_1[0][0]
dense_2 (Dense)	(None, 2048)	104448	dense_1[0][0]
reshape_1 (Reshape)	(None, 4, 4, 128)	0	dense_2[0][0]
conv2d_transpose_1 (Conv2DTranspose)	(None, 8, 8, 64)	73792	reshape_1[0][0]

```

-----
add_1 (Add)                                (None, 8, 8, 64)    0
conv2d_transpose_1[0][0]                                conv2d_2[0][0]
-----
leaky_re_lu_2 (LeakyReLU)                  (None, 8, 8, 64)    0
add_1[0][0]
-----
batch_normalization_4 (BatchNor (None, 8, 8, 64)    256
leaky_re_lu_2[0][0]
-----
conv2d_transpose_2 (Conv2DTrans (None, 16, 16, 32)    18464
batch_normalization_4[0][0]
-----
batch_normalization_5 (BatchNor (None, 16, 16, 32)    128
conv2d_transpose_2[0][0]
-----
dropout_3 (Dropout)                        (None, 16, 16, 32)    0
batch_normalization_5[0][0]
-----
conv2d_transpose_3 (Conv2DTrans (None, 32, 32, 3)    867
dropout_3[0][0]
=====
Total params: 394,549
Trainable params: 393,909
Non-trainable params: 640
-----

```

Training

```

[ ]: with device:
      autoencoder.compile(optimizer='adam', loss = "binary_crossentropy")
      history = autoencoder.fit(
          x_train_noise,
          x_train,
          batch_size = 128,
          nb_epoch = 10,
          verbose = 1,
      )

```

```

Epoch 1/10
50000/50000 [=====] - 9s 184us/step - loss: 0.5784
Epoch 2/10
50000/50000 [=====] - 7s 141us/step - loss: 0.5631
Epoch 3/10
50000/50000 [=====] - 7s 138us/step - loss: 0.5622
Epoch 4/10
50000/50000 [=====] - 7s 137us/step - loss: 0.5618
Epoch 5/10

```



```


50000/50000 [=====] - 7s 139us/step - loss: 0.5615
Epoch 6/10
50000/50000 [=====] - 7s 140us/step - loss: 0.5614
Epoch 7/10
50000/50000 [=====] - 7s 138us/step - loss: 0.5613
Epoch 8/10
50000/50000 [=====] - 7s 138us/step - loss: 0.5611
Epoch 9/10
50000/50000 [=====] - 7s 139us/step - loss: 0.5611
Epoch 10/10
50000/50000 [=====] - 7s 139us/step - loss: 0.5610

```

```

[ ]: plot_df = pd.DataFrame.from_dict({'train_loss':history.history['loss']})
plot_df.plot(logy=True, figsize=(15,5), fontsize=12)
plt.xlabel('epoch', fontsize=12)
plt.ylabel('loss', fontsize=12)
plt.show()


```



```

[ ]: # Visualization of 5 x 5 denoised images
fig = plt.figure(figsize=(32,32))
ind = 1
for i in range(1, 6, 1):
    for j in range(1, 6, 1):
        fig.add_subplot(5,5,ind)
        show(x_images[ind])
        plt.xticks([])
        plt.yticks([])
        ind+=1

```



5.4 Comparison between original, noisy and denoised images

- 1st column: original image
- 2nd column: noisy image
- 3rd column: denoised image

Images in training set

```
[ ]: # Visualization of 9 x 3 images
fig = plt.figure(figsize=(32,32))
ind = 1
for j in range(1, 10, 1):
    #-----
    # original images
```

```

fig.add_subplot(9,3,ind)
show(x_images[j])
#-----
# noisy images
fig.add_subplot(9,3,ind+1)
show(x_train_noise[j])
#-----
# denoised images
fig.add_subplot(9,3,ind+2)
show(np.squeeze(autoencoder.predict(np.expand_dims(x_images[j], 0))))
#-----
plt.xticks([])
plt.yticks([])
ind+=3

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Images in test set

```
[ ]: # Visualization of 9 x 3 images
fig = plt.figure(figsize=(32,32))
ind = 1
for j in range(1, 10, 1):
    #-----
    # original images
```

```

fig.add_subplot(9,3,ind)
show(x_test[j])
#-----
# noisy images
fig.add_subplot(9,3,ind+1)
show(x_test_noise[j])
#-----
# reconstructed images
fig.add_subplot(9,3,ind+2)
show(np.squeeze(autoencoder.predict(np.expand_dims(x_test_noise[j], 0))))
#-----
plt.xticks([])
plt.yticks([])
ind+=3

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

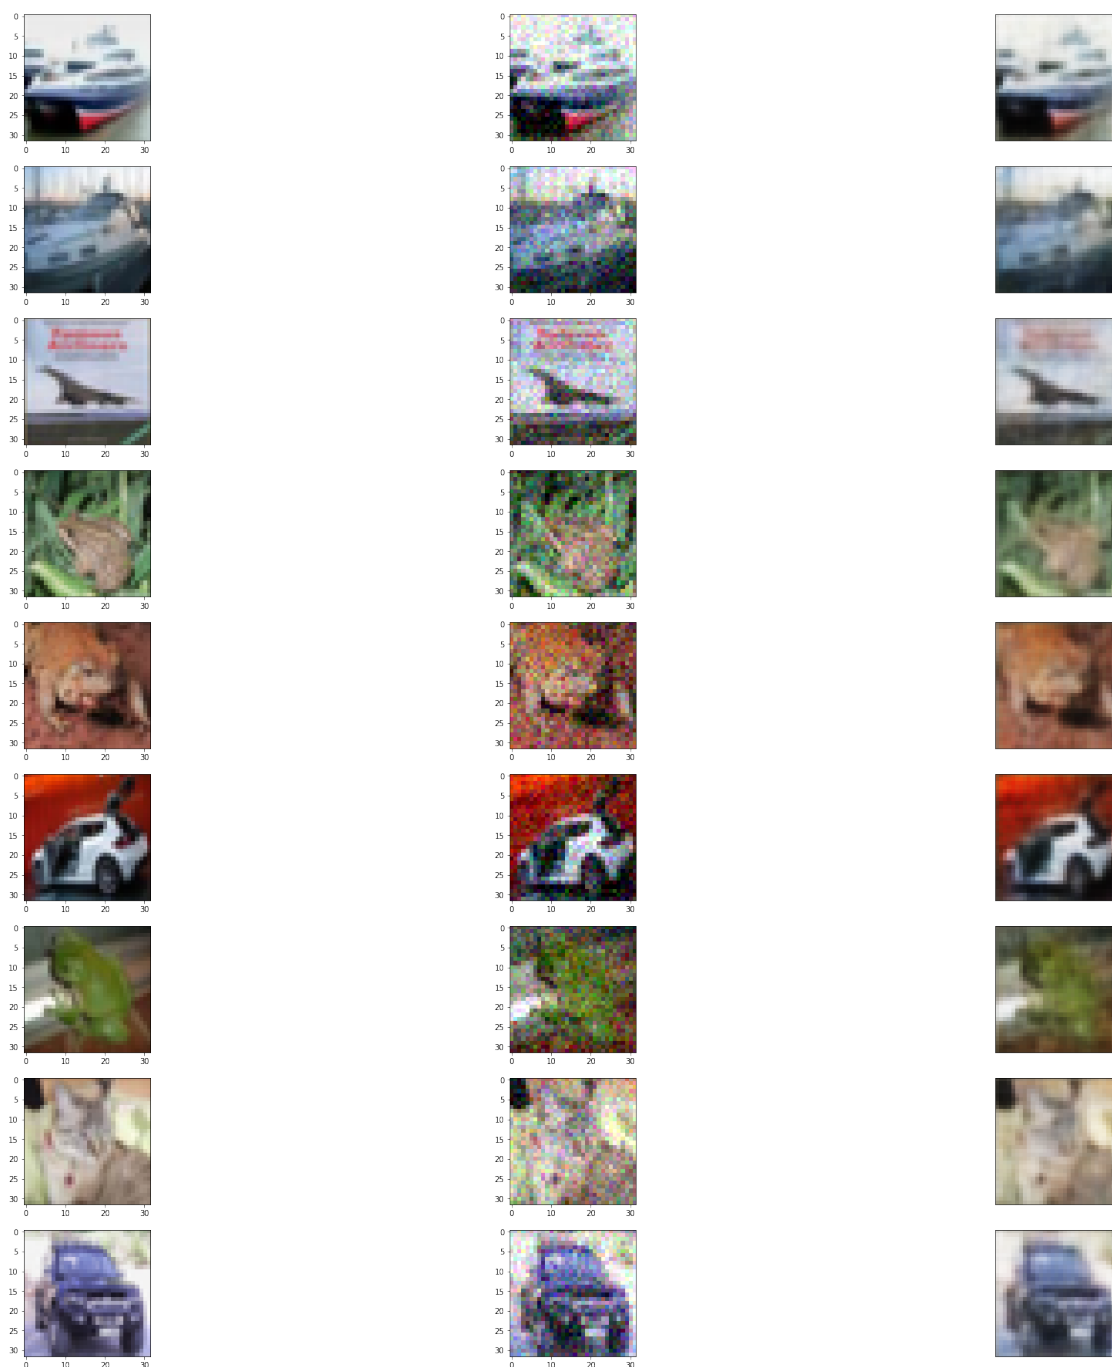
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Comments:

As we can see, the DAE (Denoising Autoencoder) is able to remove noise from images. The denoised image is somewhat distorted from the original one (as expected), but it still has quality far better than the noisy images.

End of Notebook